

# C++20 Language Features for a new Library

Alisdair Meredith  
BloombergLP

# What is this talk?

- My own *personal* vision
  - Does *not* represent ISO committee positions in any way
  - Does not represent Bloomberg official positions in any way!
- A plan to get folks with their own visions talking
  - Hope to inspire a broad vision for a new library, and supporting language, from the folks best placed to drive it

# Format

- Wednesday:  
Language Features for C++20 (and beyond)
- Thursday:  
Design pressures, pitfalls, and opportunities for std2
- Friday:  
“Workshop” is a feedback session to try to inspire a collective vision paper for ISO in July

# Inspiration

- First attempt to constrain standard library with concepts, around a decade ago
- Failed for several reasons
  - Concepts were not ready, and becoming more complex by patching them late
  - Library insisted on complete compatibility with existing code
    - produced too many concepts
    - deprived concepts of power, by adding concepts till nothing broke
- Concern that a concept-based library would have to break backwards compatibility, or compromise itself so badly it may not be worth the effort

# Why Now?

- C++17 reserved a family of namespaces for future libraries
- We have long dreamed of a future incompatible library, fixing many defects
- Concepts are coming, and that seems the likely spark to start such an overwhelming project

# Why *this* talk?

- New library should have solid foundations
- What are the missing or anticipated language features that should be integral to any new library design?
- Can we identify those that are critical, and ensure they land in C++20?

# How to pick features

- Language feature will change the nature of abstraction in the language
- Language feature would have a clear impact on library interfaces
- Language feature would have a fundamental impact on library organization

# Format of this talk

- Assume a broad familiarity with the ideas behind many of the proposals
- Will zoom in on a couple of interesting that merit further discussion as part of standardization
- Will not provide in depth tutorials - each might be a session in its own right
- Try to bring focus back to library design



# Modules

# Modules

- The single, most fundamental change
- Affects every library entity, whether a class, function, template, or alias
  - Apart from (maybe) macros
- Also the smallest, as more about packaging than content

# State of the art

- Clang has been evolving their model of modules for around 5 years
- MS have been shipping their experimental modules since VC2015, and upgrade in subsequent releases
- ISO close to finalizing a modules TS, drawing on experience from both schools of thought
- gcc now looking to implement modules TS

# What is a module?

- Library level abstraction
- A collection of multiple TUs that can be accessed through a single import directive
- C++ code, not repeated textual copy/paste of `#include`
- Ideally one trip to file system to access file, not hundreds of chained headers per TU

# Basic syntax

- `module identifier; // names the module`
- `export namespace blah {  
 // exported declarations, and class definitions  
}`
- `import identifier; // access contents of the  
 // named module`

# Module contents

- a single interface file declares all exports
- definitions for all exported contents, potentially split across multiple TUs
- Additional declarations and definitions are not exported, so not accessible outside the module
  - truly private, not subject to ODR *across modules*
  - Obviously ODR applies within a module

# Module Constraints

- No cycles between modules
- Import graph must be a DAG  
*Directed Acyclic Graph*
- Mutually dependent classes must reside in the same module, but maybe different TUs
- Templates may circumvent these rules, but retain the import DAG

# How to split a module?

- How do multiple TUs access the class definitions from their interface file?
  - reinventing `#include`?
  - more syntax needed?
- build system recognizes non-interface TUs, and implicitly imports the interface for every TU with that module identifier?



# Open questions: 1

- Can a module export a macro?
  - *My* preference: no  
modules export code, macros are preprocessor utilities that manipulate source text  
access macros using preprocessor, with a traditional `#include`
- Transition to from C++17 to modules may require mixed mode support, including macros

# Open Questions: 2

- Must the module declaration be the first substantive line in a TU?

# Open Questions: 3

- templates?!
- how do we avoid recreating the problems of C++98 exported templates?

# Bloomberg experience

- Bloomberg architecture:
  - A component is a .h/.cpp pair
  - A package comprises multiple components
  - A package group comprises multiple packages
  - No higher level of aggregation
  - No cycles between components

# Bloomberg Modules?

- A component forms a single module
- A package is a module comprising component modules
- A package group is a module comprising package modules (never components directly)

# Early Experience

- Package modules export a *dependency* on component modules
  - Modules are always assembled by-reference, never by-value
- Result is we would distribute thousands of modules - granularity feels wrong
- Suggestion this might be addressed by a future module partitions feature

Concepts

# Concepts

- Major feature, long time coming
- No checked definitions (and that is a good thing)
- Ranges library as pilot proof-of-concept(s)
- Terse syntax essential for lambdas
- Forwarding problem with terse syntax



# Assumptions

- This is a technically literate audience that is aware, in a broad sense, that concepts is a language feature in development for over a decade
- Understand the basic idea of a syntax to constrain parameters of generic code
- Serve as documentation for both the user and the compiler

# Current State of the Art

- Concepts TS published November 15, 2015
  - No longer Concepts *Lite*
- gcc 6.1 shipped a supported implementation April 2016
- L(E)WG hard at work building on this with Ranges TS as foundation on library work

# Features of Note

- Concepts can be declared with either a variable or a function syntax
- Concepts always return `bool`, or are `constexpr bool` variables, but must still state the `bool`
- `requires` expressions are usable only in `requires` clauses
  - relaxed since TS was published

# Features of Note

- A much more expressive and efficient syntax for SFINAE conditions
- Constraint-based overloading on otherwise identical signatures e.g., overloading for forward iterators and random access iterators
- Extremely flexible syntax for declaring constraints

# And there's more...

- “normal” form uses concepts like types and does away with the template head
- `auto` can now be used for a function parameter to create an unconstrained template
- Would this be better handled by an `Auto` library concept?

# A quick example

```
void foo(ArgType&& arg) {  
    bar(arg);  
}
```

# A quick example

Q: Can we call foo with an lvalue?

```
void foo(ArgType&& arg) {  
    bar(arg);  
}
```

# A quick example

Q: Can we call foo with an lvalue?

```
void foo(ArgType&& arg) {  
    bar(arg);  
}
```

A: It depends...



# A quick example

Q: Can we call foo with an lvalue?

```
void foo(ArgType&& arg) {  
    bar(arg);  
}
```

A: It depends...

Is **ArgType** a type or a concept?

# Possible Solution?

Real syntax for forwarding references

```
void foo(ArgType&&& arg) {  
    bar(arg);  
}
```

- `arg` behaves strictly as an lvalue or an rvalue, depending on call context
- Whether `ArgType` is a type or a concept
- Essentially a mini-template when `ArgType` is a type
- We need to find a better syntax!

# Possible Solution?

Real syntax for forwarding references

```
void foo(ArgType&&& arg) {  
    bar(arg);  
}
```

- Eliminates need for `std::forward`?
  - Unconstrained templates retain, but deprecate, old behavior
- `&&` is always an rvalue unless template keyword is present
  - not an easy thing to express today without SFINAE hacks

# Multiple Deductions of a Concept

```
template < typename ExecutionPolicy
          , typename ForwardIterator1
          , typename ForwardIterator2
          >
void copy( ExecutionPolicy&& exec
          , ForwardIterator1 first
          , ForwardIterator1 last
          , ForwardIterator2 result
          );
```

# Multiple Deductions of a Concept

```
template < ExecPol  ExecutionPolicy
          , FwdIter  ForwardIterator1
          , FwdIter  ForwardIterator2
        >
void copy( ExecutionPolicy&& exec
          , ForwardIterator1  first
          , ForwardIterator1  last
          , ForwardIterator2  result
        );
```

# Multiple Deductions of a Concept

Do `first`, `last`, and `result` deduce to the same or different types?

```
void copy( ExecPol&& exec
          , FwdIter   first
          , FwdIter   last
          , FwdIter   result
          );
```

# Multiple Deductions of a Concept

I propose adding a tag to guide deduction

```
void copy( ExecPol&& exec
           , FwdIter.1 first
           , FwdIter.1 last
           , FwdIter.2 result
           );
```

# Multiple Deductions of a Concept

I propose adding a tag to guide deduction

```
void copy( ExecPol&& exec
           , FwdIter.1 first
           , FwdIter.1 last
           , FwdIter.2 result
           );
```

Without a tag, all named concepts deduce to the same type, and so behave on the page just like a type does today.



# Tagged Concepts

For consistency, tag would apply through whole deduced scope, e.g., for variable declarations

```
void func(ArgType.1 x, ArgType.2 y) {  
    ArgType.1  z{x};    // copy x  
    ArgType.2  *p{&y};  // take address of y  
    ArgType.1  q{y};    // ...  
}
```

Contracts

# Contracts

- Run-time analog of concepts
- Support better testing/support for software
- Support static analysis tools
- Clearly specify the rights and responsibilities of the caller of a function

# Basics

- Proposal uses modified attribute syntax
- `[[expects : predicate]]` indicate preconditions the caller must satisfy
- `[[ensures : predicate]]` gives a guarantee from the function if it returns
- `[[assert : predicate]]` can be used in a function body to check preconditions that are in the doc.

# Violation Handlers

- Predicates may be checked at runtime, and a handler will be called if the predicate is false
- Specify (on the build line) the handler to respond to a failed check
  - defaults to calling `abort`
  - if a handler throws out of a `noexcept` function, it violates the exception specification, and terminates
- Extra flag to allow handlers to return back to caller
  - e.g., to install a logging handler when annotating pre-existing code

# Restrictions

- Handler cannot be changed at runtime
  - although could install a handler with its own registration scheme
- predicates respect the access specifier of the function
  - public functions can check only against public members
  - protected functions cannot check against private members

# Checking Levels

- Not all contract checks are cheap
  - e.g., post-condition that a sort function produces a permutation of the original sequence
- Allow annotation to express a simple cost
  - default : always check
  - audit : check only if command line enables expensive checks, typically violating complexity guarantees
  - axiom : never run, but useful information, e.g., for static analysers, `is_valid_pointer()`
- Allow the whole feature to be disabled on the command line

# Concerns

- What information should be passed to the handler?
  - Should we have options to restrict detail to save program size?
- Are we concerned about violation handlers racing?
  - (should not be messing with shared state anyway)
- Should library document checking strength?
- Are highly detailed contracts on function signatures helpful or harmful?
  - How much library doc should move into code?



# Quick Example

Could this replace `at()` in a new library?

Should we mark the functions `noexcept`?

```
template <typename T, typename Alloc>
struct vector {
    T & operator[](size_type index)
        [[expects : index < size()]];

    T & insert(T const &)
        [[ensures : !empty()]];
};
```

# Coroutines

# Coroutines

- Championed by Gor Nishanov, with a TS currently under ballot in ISO
- Going TS route, rather than straight-to-standard, to allow time for a competing proposal from Google/ASIO folks

# Coroutines

- Allow you to exit/return from a function, and then resume where you left off
- enables *generators* : ranges defined to execute a function
- Asynchronous coroutines return a future-like type

# Language Impact

- 3 new keywords
  - `co_await`
  - `co_return`
  - `co_yield`
- Allow `co_await` on range-based for loops
- Coroutine traits allow for custom implementations

# Why coroutines are important

- Simplify use of asynchronous callbacks in concurrent code
- Library interface should assume feeding generators into algorithms, much like we might expect iterator adaptors to avoid many variations in C++98 algorithms

# Quick Example

```
generator<int> coro(int from, int to) {  
    while (from != to) {  
        co_yield from++;  
    }  
}  
  
int main() {  
    for (int x : coro(3, 7)) {  
        cout << x << endl;  
    }  
}
```

# Second example

```
async_stream<int> Ticks() {  
    for (int tick = 0; ; ++tick) {  
        co_yield tick;  
        co_await 1ms;  
    }  
}
```

```
future<int> Sum(async_stream<int> & input) {  
    int sum = 0;  
    for co_await(auto&& v : input) {  
        sum += v;  
    }  
    co_return sum;  
}
```



# Unified Call Syntax

# Unified Call Syntax

- There is a split in preferred convention between libraries preferring member functions over free-functions accessing the public API of a class
- std library is split on this, duplicating more members as free functions to be called via ADL
- It would be nice if the language would transparently allow both free and member-syntax invocation for the same function
  - Simpler to write generic code without requiring a function to be a member (or not)

# Concerns

- Managing the overload set of members/non-members without introducing ambiguities or other surprises
  - One suggestion to use a new `.call(x,y)` syntax to invoke the uniform rules
- Does every member name become an extension point, making generic code nigh impossible to write in the abstract?

# Why it matters?

- Do we want to carry on extending the idiom of `begin/end/cbegin/cend/rbegin/rend/crbegin/crend/size/empty/data`?
- Issue matters less if we know uniform syntax available

# Default Comparisons

# Comparison Operators

- Long standing desire to generate, or default, comparison operators
- `std::relops` has not been a popular choice for this, over the years
  - we are lacking a tag-type to derive from, which would drive ADL
- Can simplify implementation by using `std::tie` on members, but fragile if class evolves.

# Spaceship Operator

- Proposed by Herb Sutter at most recent Kona meeting, and well received
- Basic idea: `operator<=>` returns -ve, 0, +ve to indicate less-than, equal-to/equivalent, or greater-than

# Devil is in the Details

- Not all types support ordering, so want boolean result if types are only equality comparable
- May need to know if types are strictly ordered, or only partially ordered
- Result of operator is not an int, but a tag-type wrapping an int, that encodes this behavior



# Broader Semantics

- If a type defines only `operator<=>`, the other operators can be synthesized on demand
- Demand a consistent definition

<code>a == b</code>	<code>a &lt;=&gt; b == 0</code>	<code>0 == b &lt;=&gt; a</code>
<code>a != b</code>	<code>a &lt;=&gt; b != 0</code>	<code>0 != b &lt;=&gt; a</code>
<code>a &lt; b</code>	<code>a &lt;=&gt; b &lt; 0</code>	<code>0 &lt; b &lt;=&gt; a</code>
<code>a &lt;= b</code>	<code>a &lt;=&gt; b &lt;= 0</code>	<code>0 &lt;= b &lt;=&gt; a</code>
<code>a &gt; b</code>	<code>a &lt;=&gt; b &gt; 0</code>	<code>0 &gt; b &lt;=&gt; a</code>
<code>a &gt;= b</code>	<code>a &lt;=&gt; b &gt;= 0</code>	<code>0 &gt;= b &lt;=&gt; a</code>

# Example Implementation

Example taken directly from P0515R0

```
std::strong_ordering operator <=> (const TotallyOrdered& that) const {  
    if (auto cmp = (Base&)(*this) <=> (Base&)that;    cmp != 0) return cmp;  
    if (auto cmp = last_name. <=> that.last_name;    cmp != 0) return cmp;  
    if (auto cmp = first_name <=> that.first_name;    cmp != 0) return cmp;  
    return tax_id <=> that.tax_id;  
}
```

# Example Implementation

Example taken directly from P0515R0

```
std::strong_ordering operator <=> (const TotallyOrdered& that) const {  
    if (auto cmp = (Base&)(*this) <=> (Base&)that; cmp != 0) return cmp;  
    if (auto cmp = last_name. <=> that.last_name; cmp != 0) return cmp;  
    if (auto cmp = first_name <=> that.first_name; cmp != 0) return cmp;  
    return tax_id <=> that.tax_id;  
}
```

Note neat use of C++17 `if` with a variable declaration

# Default Definition?

- Request a default definition with `= default`
- Like existing special members, the exception specification is also deduced for defaulted definitions

Overloading  
constexpr

# constexpr overloads

- A real desire to make more of the library `constexpr`, or at least `constexpr`-friendly
- don't want to pay a cost for less efficient run-time behavior, restricted to `constexpr`-friendly syntax
  - e.g., `asm`, or use of intrinsics for `memcpy` etc.

# constexpr overloads

- solution 1:  
allow overloading functions on `constexpr`
- solution 2:  
introduce an `is_constexpr_eval` intrinsic that  
can be used as predate for `if constexpr`
- and perhaps assume all functions may be  
`constexpr`, and fail only if not usable at call-  
point

# Quick example

```
template <>
constexpr
char * char_traits<char>::copy( char * s1
                                , char const * s2
                                , size_t n)
{
    if constexpr (is_constexpr_eval) {
        while (n--) { *s1++ = *s2++; }
        return s1;
    }
    else {
        return strcpy(s1, s2, n);
    }
}
```



Reflection

# Reflection

- Reflection is a popular facility in many languages with a managed runtime
- C++ libraries emulate compile-time reflection with SFINAE tricks, some of which is obviated by concepts
- Runtime reflection would (probably) add considerable cost to object sizes, and may start to dictate more implementation details for language representation than compiler vendors are happy with

# Why do we want it?

- Adapting software on the fly at compile and runtime
  - e.g., creating data bindings after querying a database for its schema at runtime
  - e.g., querying a class for its members and signatures, to drive an automated test driver
  - etc.

# Progress

- Reflections study group is broadening scope to cover meta-programming in general, as a typical participant in compile-time reflection schemes
- First results will be reviewed in EWG over the next few meetings...

# Transactional Memory

# Transactional Memory

- An active research topic to produce a simpler concurrent programming model
- Modeled on database transactions
  - optimistically run some code, and commit to memory if no races are detected
  - otherwise roll back, and try again
- Transaction-safe code must be explicitly marked
  - Library must document/mark transaction safety

# Transactional Memory

- Expected to have a minor cost with hardware support, larger impact with software solution
  - Hardware support has been coming for some time
- We have an experimental TS
  - Implementation available since gcc 6.1
- Expecting further work

# Proxy Support



# Proxy Support

- how can we better support iterators returning proxies?
- should we be able to overload operator dot?
  - can we write generic code when both `operator .` and `operator &` may be overloaded?
- Nervous about `addressof` all over again - how do I get address out of my proxy?
- (not possible for true rvalues)

# Proxy Support

- I got scared and stopped looking at this point
- Consider me a source of FUD!

# What do we want?

- Looking ahead to next standards, which features are critical for a new library
  - Which features did I miss?
- Push for those in C++20
  - get involved
- Which features would redirect significant effort if they arrived later?
- Which features can we absorb incrementally in C++23 and beyond?

# My Priorities

- Modules
- Concepts
- Contracts
- Coroutines

# My Wish List

- Runtime Reflection
- My own enhancements to concepts ;)
- Allocator support (see tomorrow!)
- Transactional memory

# Nice to Have

- Spaceship operator
- Better control of constexpr code generation

# Would rather rule out

- Universal calling syntax
- Overloading operator dot